

# Introduction to Docker

Peng Xiao

GitHub: xiaopeng163

Network Consulting Engineer

Cisco Systems

# Contents

---

- The Metrix from Hell
- What is Container/Docker?
- Docker Basic Usage
- Docker Networking Deep Dive
- Docker Compose
- Docker Swarm
- Ecosystem and Standardization

# The Challenge

Multiplicity of Stacks



Static website

nginx 1.5 + modsecurity + openssl + bootstrap 2



Background workers

Python 3.0 + celery + pyredis + libcurl + ffmpeg + libopencv + nodejs + phantomjs



User DB

postgresql + pgvll + vll



Queue

Redis + redis-sentinel



Analytics DB

hadoop + hive + thrift + OpenJDK



Web frontend

Ruby + Rails + sass + Unicorn



API endpoint

Python 2.7 + Flask + pyredis + celery + psycopg + postgresql-client

Do services and apps  
interact  
appropriately?



Multiplicity of  
hardware  
environments



Development VM



QA server

Customer Data Center



Public Cloud



Production Cluster



Disaster recovery














Production Servers

Contributor's laptop



Can I migrate  
smoothly and  
quickly?

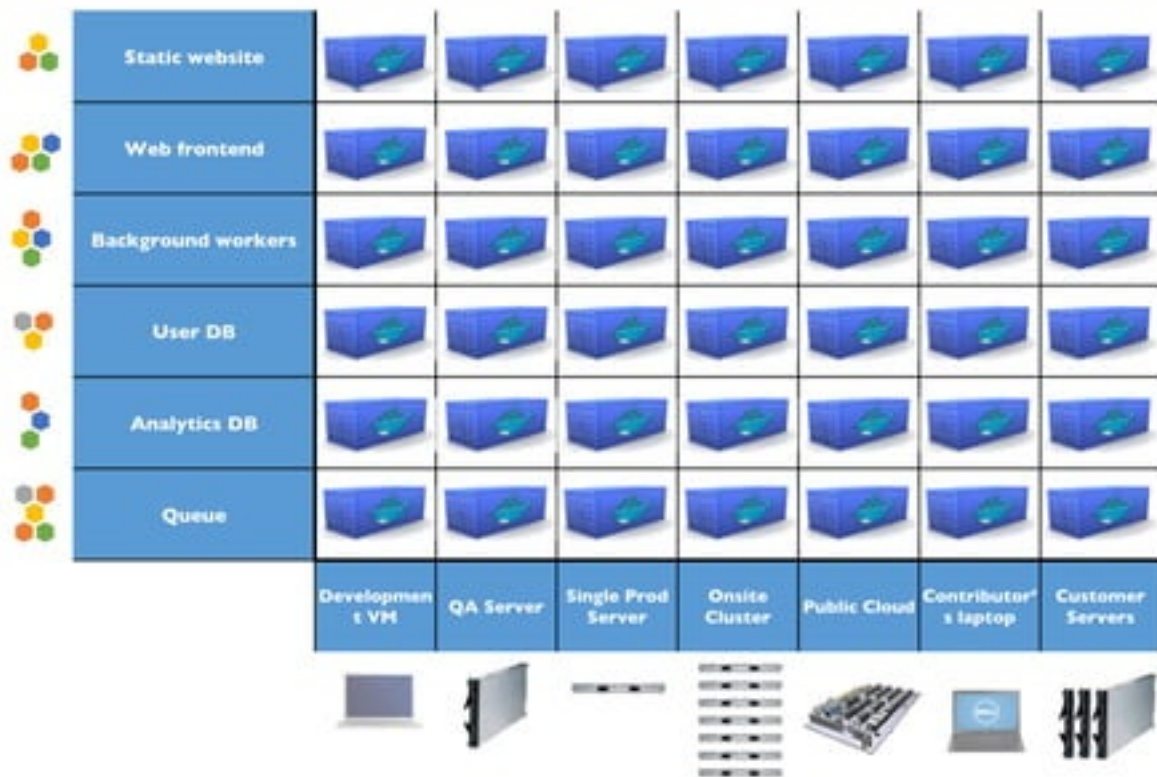
# The Matrix From Hell

	Static website	?	?	?	?	?	?	?
	Web frontend	?	?	?	?	?	?	?
	Background workers	?	?	?	?	?	?	?
	User DB	?	?	?	?	?	?	?
	Analytics DB	?	?	?	?	?	?	?
	Queue	?	?	?	?	?	?	?
		Development VM	QA Server	Single Prod Server	Onsite Cluster	Public Cloud	Contributor's laptop	Customer Servers
								

# Docker is a shipping container system for code



# Docker eliminates the matrix from Hell





# Why Developers Care

---

- Build once...(finally) run anywhere\*
  - A clean, safe, hygienic and portable runtime environment for your app.
  - No worries about missing dependencies, packages and other pain points during subsequent deployments.
  - Run each app in its own isolated container, so you can run various versions of libraries and other dependencies for each app without worrying
  - Automate testing, integration, packaging...anything you can script
  - Reduce/eliminate concerns about compatibility on different platforms, either your own or your customers.
  - Cheap, zero-penalty containers to deploy services? A VM without the overhead of a VM? Instant replay and reset of image snapshots? That's the power of Docker

\* With the 0.7 release, we support any x86 server running a modern Linux kernel (3.2+ generally. 2.6.32+ for RHEL 6.5+, Fedora, & related)



# Why Devops Cares?

---

- Configure once...run anything
  - Make the entire lifecycle more efficient, consistent, and repeatable
  - Increase the quality of code produced by developers.
  - Eliminate inconsistencies between development, test, production, and customer environments
  - Support segregation of duties
  - Significantly improves the speed and reliability of continuous deployment and continuous integration systems
  - Because the containers are so lightweight, address significant performance, costs, deployment, and portability issues normally associated with VMs



# Contents

---

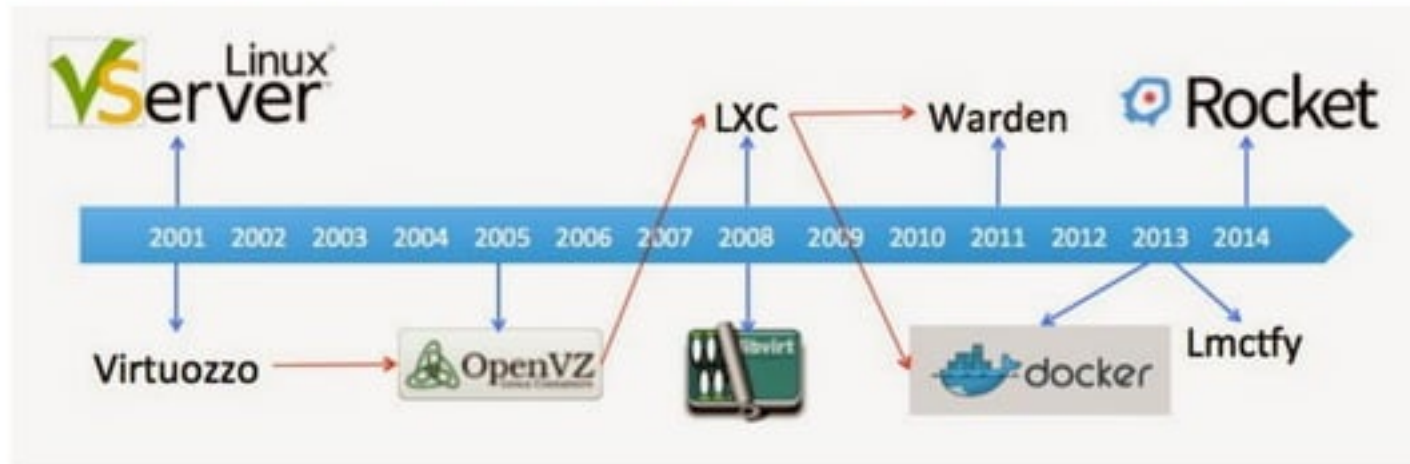
- The Metrix from Hell

- **What is Container/Docker?**

- Docker Basic Usage
- Docker Networking Deep Dive
- Docker Compose
- Docker Swarm
- Ecosystem and Standardization

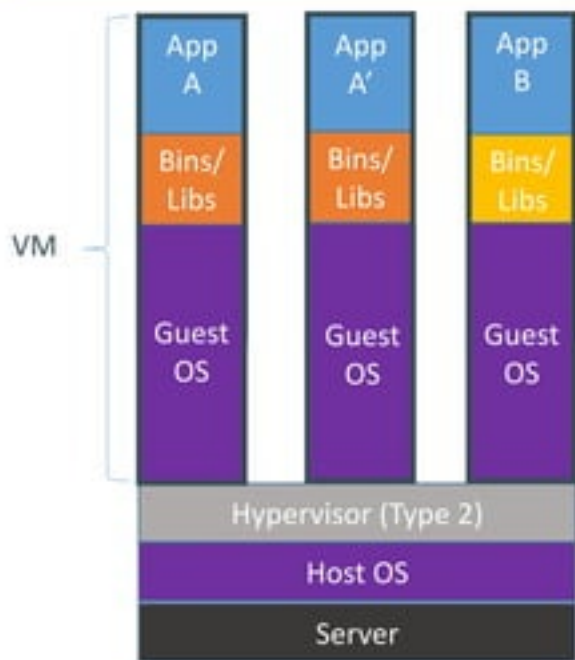
# Container History

*Docker is nothing new, it just show in the right time*



<http://kiwenlau.blogspot.com/2015/01/linux-container-technology-overview.html>

# Containers vs. VMs

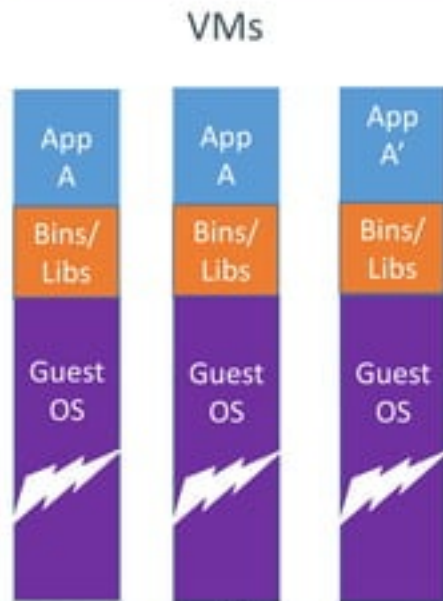


Containers are isolated, but share OS and, where appropriate, bins/libraries

...result is significantly faster deployment, much less overhead, easier migration, faster restart



# Why are Containers lightweight?



VMs

Every app, every copy of an app, and every slight modification of the app requires a new virtual server

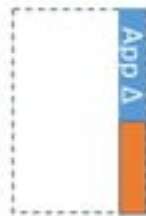
## Containers



Original App  
(No OS to take up space, resources, or require restart)



Copy of App  
No OS. Can Share bins/libs



Modified App

Copy on write capabilities allow us to only save the diffs Between container A and container A'

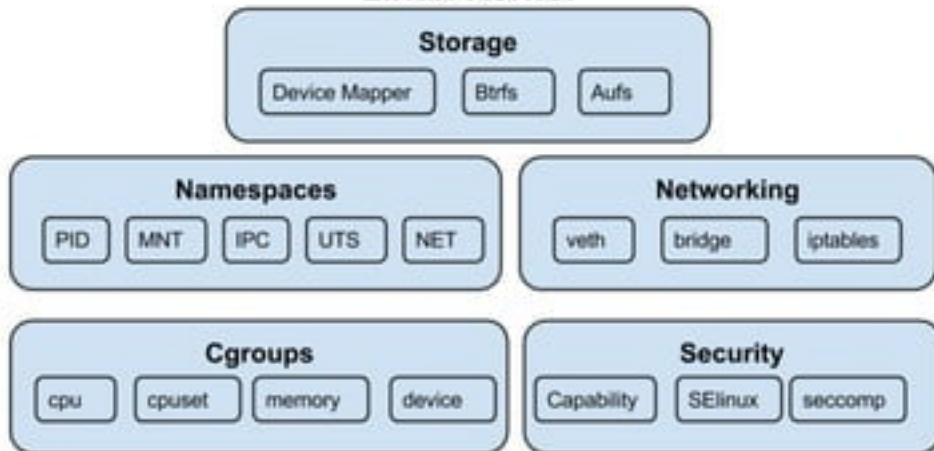
# Docker Inside

---



---

## Linux Kernel



## Docker Inside (Deep Dive)

---

- A Deep Dive Into Linux Containers For Engineers Interested In The Gritty Details. <http://docker-saigon.github.io/post/Docker-Internals/>
- Cgroups, namespaces, and beyond: what are containers made from? (DockerCon Europe 2015) <https://goo.gl/25KtpZ>



# Contents

---

- The Metrix from Hell
- What is Container/Docker?
- **Docker Basic Usage**
- Docker Networking Deep Dive
- Docker Compose
- Docker Swarm
- Ecosystem and Standardization

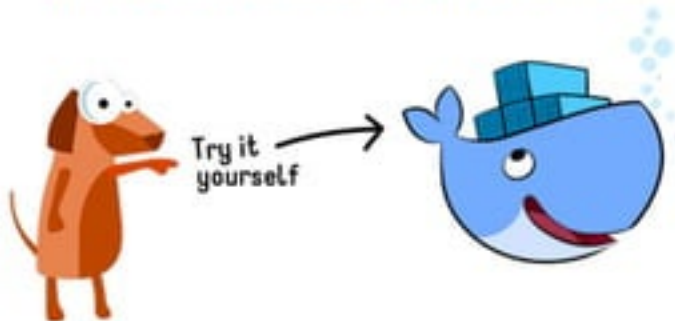




# Docker Engine

---

- Docker Daemon + REST(ish) API
- Docker Client (CLI) talks to the API
- Daemon manages the Docker containers
- Start it with: `docker -d`
- `docker version` to test if docker is setup correctly
- Docker Engine (Linux, Mac, Windows)
  - <https://docs.docker.com/engine/installation/>



# Docker Toolbox

- Docker Toolbox is an installer for quick setup and launch of a Docker environment on **older Mac and Windows systems** that do not meet the requirements of the new Docker for Mac and Docker for Windows apps.



# Docker Machine

- Machine makes it really easy to create Docker hosts on your computer, on cloud providers and inside your own data center. It creates servers, installs Docker on them, then configures the Docker client to talk to them.
- Drivers exist for:
  - AWS
  - DigitalOcean
  - Azure
  - Google Compute Engine
  - Rackspace
  - OpenStack
  - **Virtualbox**
  - **VMWare Fusion**
  - VMWare vSphere
  - Hyperv

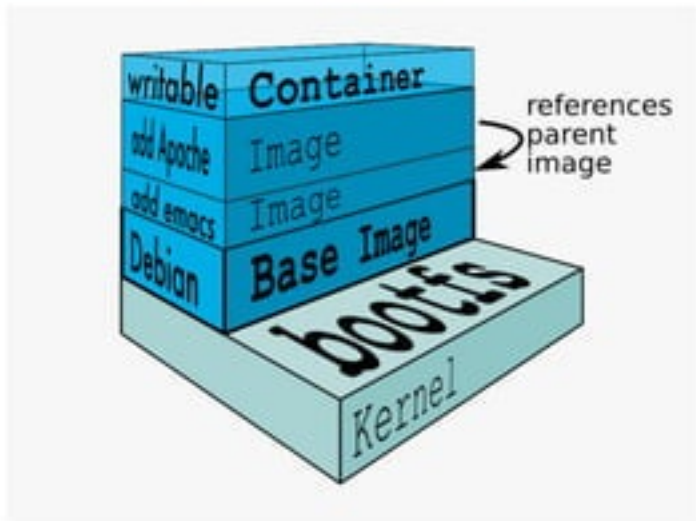
## Docker Machine



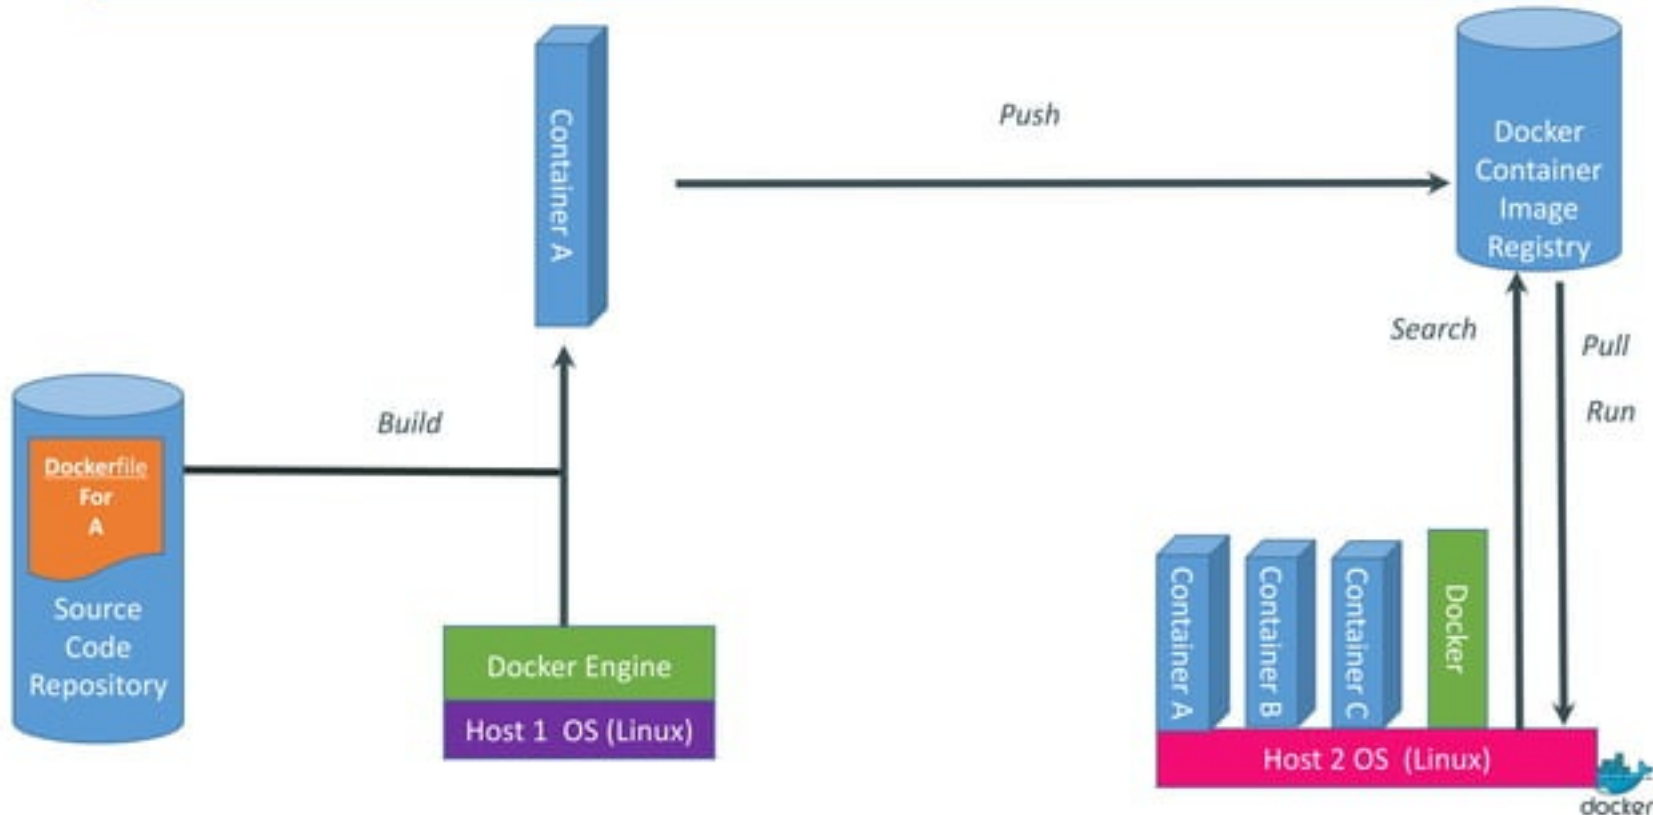
# Docker Images vs Containers

---

- Images: About storing or moving your app
- Containers: About running your app



# Docker Workflow—Build/Pull/Push



# Docker Image

- Download, list, remove images

```
$ docker pull ubuntu:14.04
14.04: Pulling from library/ubuntu

04cf3f0e25b6: Pull complete
d5b45e963ba0: Pull complete
a5c78fda4e14: Pull complete
193d4969ca79: Pull complete
d709551f9630: Pull complete
Digest: sha256:edb984703bd3e0981ff541a5b9297ca1b81fde6e6e8094d86e390a38ebc30b4d
Status: Downloaded newer image for ubuntu:14.04
```

```
$ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
ubuntu	14.04	aee2b63c4946	12 hours ago	187.9 MB

```
$ docker rmi aee2b63c4946
Untagged: ubuntu:14.04
Deleted: sha256:aee2b63c49461fcae4962e4a8043f66acf8e3af7e62f5ebceb70b181d8ca01e0
Deleted: sha256:50a2a0443efd0936b13eebb86f52b85551ad7883e093ba0b5bad14fec6ccf2ee
Deleted: sha256:9f0ca687b5937f9ac2c9675065b2daf1a6592e8a1e96bce9de46e94f70fbf418
Deleted: sha256:6e85e9fb34e94d299bb156252c89dfb4dcec65deca5e2471f7e8ba206eba8f8d
Deleted: sha256:cc4264e967e293d5cc16e5def86a0b3160b7a3d09e7a458f781326cd2cedb1
Deleted: sha256:3181634137c4df95685d73bfbc029c47f6b37eb8a80e74f82e01cd746d0b4b66
```



<https://hub.docker.com/>

# Dockerfile

---

## Create a Dockerfile

```
$ more Dockerfile
FROM      ubuntu:14.04
MAINTAINER xiaoquwl@gmail.com
RUN       apt-get update && apt-get install -y redis-server
EXPOSE    6379
ENTRYPOINT ["/usr/bin/redis-server"]
```

## Build a image

```
$ docker build -t xiaopeng163/redis:0.1 .
$ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
xiaopeng163/redis	0.1	ccbca61a8ed4	7 seconds ago	212.4 MB
ubuntu	14.04	3f755ca42730	2 days ago	187.9 MB



# Docker Container

- Create and start a container

```
$ docker run -d --name demo xiaopeng163/redis:0.1
4791db4ff0ef5a1ad9ff7c405bd7705d95779b2e9209967fffbef66cbaee80f3a
```

```
$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
4791db4ff0ef	xiaopeng163/redis:0.1	"/usr/bin/redis-serve"	5 seconds ago	Up 4 seconds	6379/tcp	demo

- Inside of the container

```
$ docker exec demo ps -ef
```

UID	PID	PPID	C	STIME	TTY	TIME	CMD
root	1	0	0	06:43	?	00:00:00	/usr/bin/redis-server *:6379
root	10	0	0	06:59	?	00:00:00	ps -ef

```
$ docker exec -it demo bash
```

```
root@4791db4ff0ef:/# ps -ef
```

UID	PID	PPID	C	STIME	TTY	TIME	CMD
root	1	0	0	06:43	?	00:00:00	/usr/bin/redis-server *:6379
root	13	0	0	07:00	?	00:00:00	bash
root	17	0	0	07:00	?	00:00:00	bash
root	30	17	0	07:00	?	00:00:00	ps -ef

# Docker Container

- Stop and remove a container

```
$ docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED             STATUS              PORTS              NAMES
4791db4ff0ef       xiaopeng163/redis:0.1  "/usr/bin/redis-serve"  20 minutes ago      Up 20 minutes      6379/tcp           demo
$ docker stop demo
demo
$ docker ps -a
CONTAINER ID        IMAGE               COMMAND             CREATED             STATUS              PORTS              NAMES
4791db4ff0ef       xiaopeng163/redis:0.1  "/usr/bin/redis-serve"  20 minutes ago      Exited (0) 7 seconds ago      demo
$ docker rm demo
demo
$ docker ps -a
CONTAINER ID        IMAGE               COMMAND             CREATED             STATUS              PORTS              NAMES
```

<http://docker-k8s-lab.readthedocs.io/en/latest/docker/docker-cli.html>

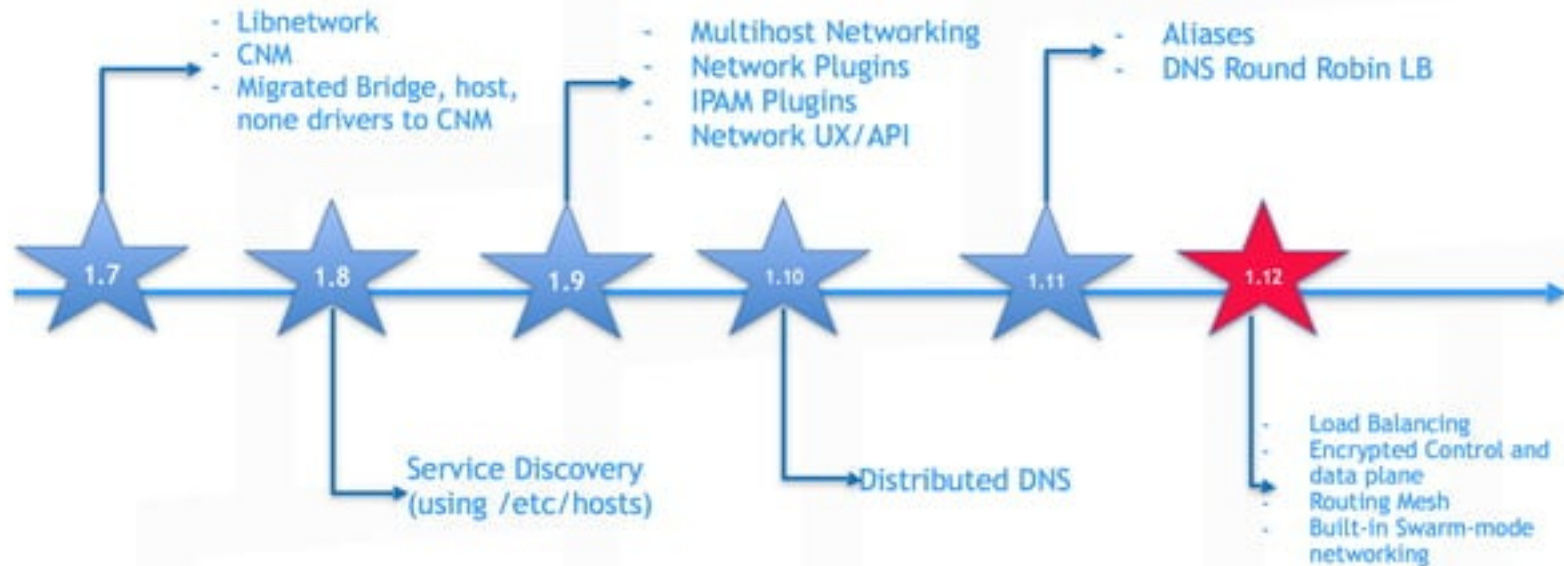
# Contents

---

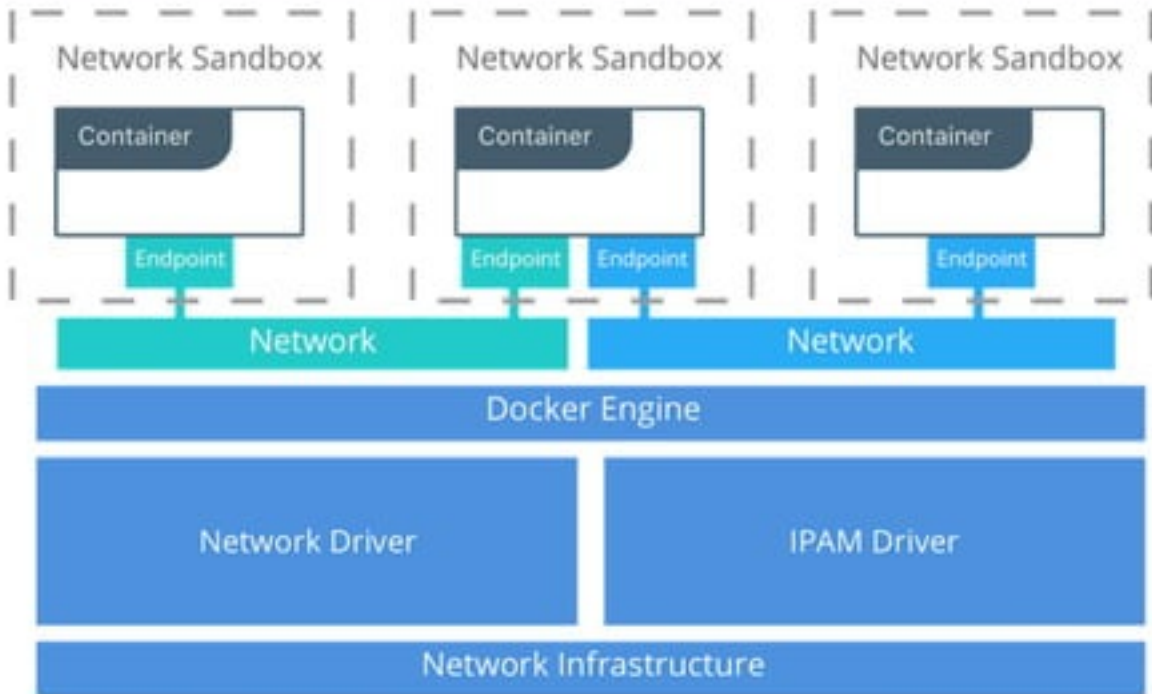
- The Metrix from Hell
- What is Container/Docker?
- Docker Basic Usage
- **Docker Networking Deep Dive**
- Docker Compose
- Docker Swarm
- Ecosystem and Standardization



# Docker Networking

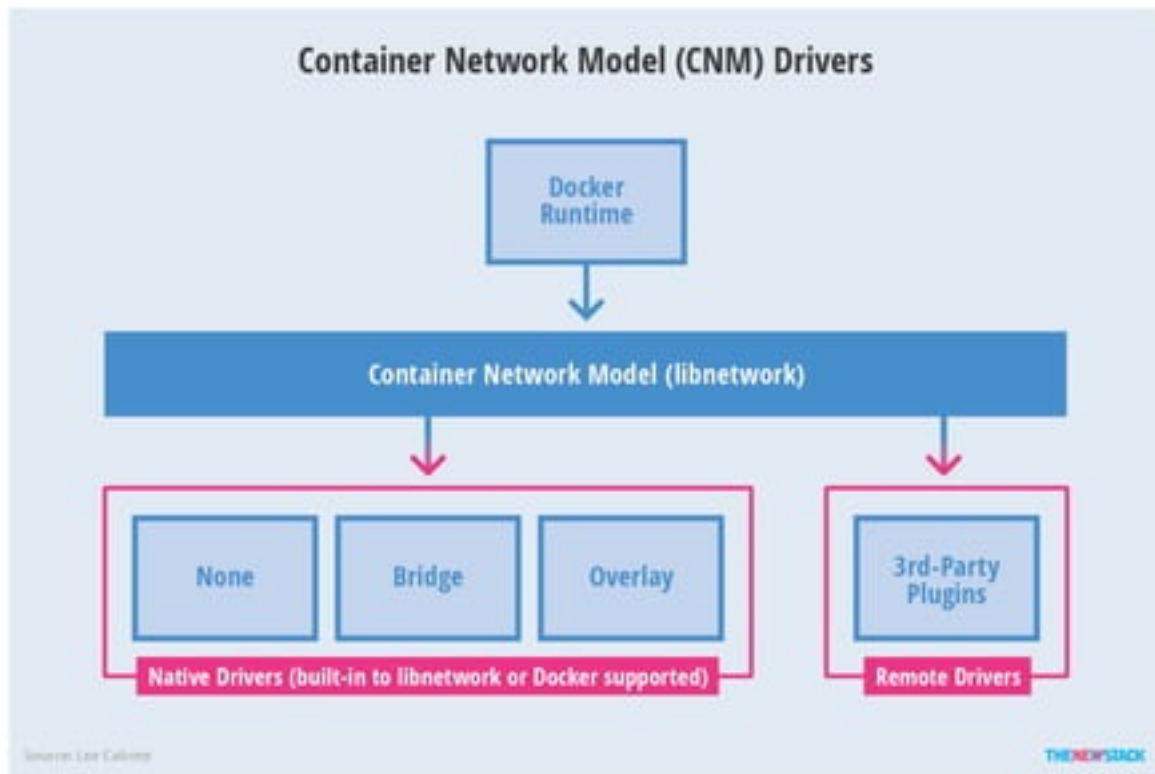


# Container Network Model



<https://github.com/docker/libnetwork/blob/master/docs/design.md>

# CNM Drivers



## Demo: Docker Bridge Network

---

- <http://docker-k8s-lab.readthedocs.io/en/latest/docker/bridged-network.html>





# Docker Multi-Host Networking

---

- Tunnel
  - Docker build-in overlay network: VXLAN
  - OVS: VXLAN or GRE
  - Flannel: VXLAN or UDP
  - Weave: VXLAN or UDP
- Routing
  - Calico: Layer 3 routing based on BGP
  - Contiv: Layer 3 routing based on BGP

<http://blog.dataman-inc.com/shurenyun-docker-133/>

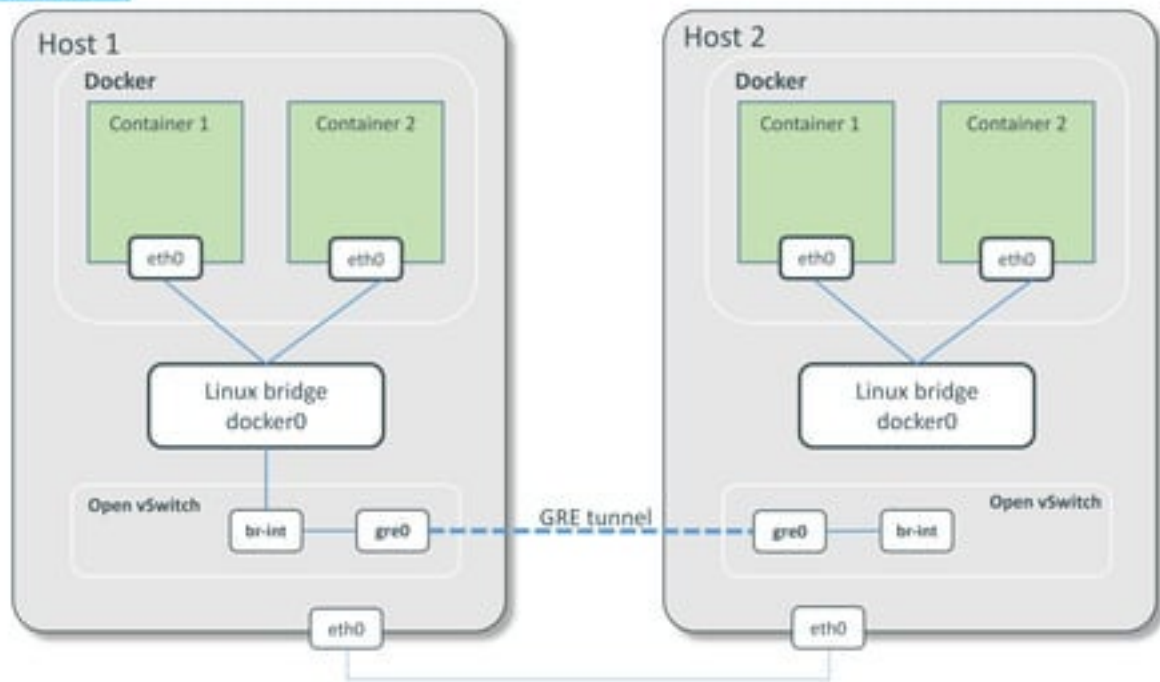


Contiv



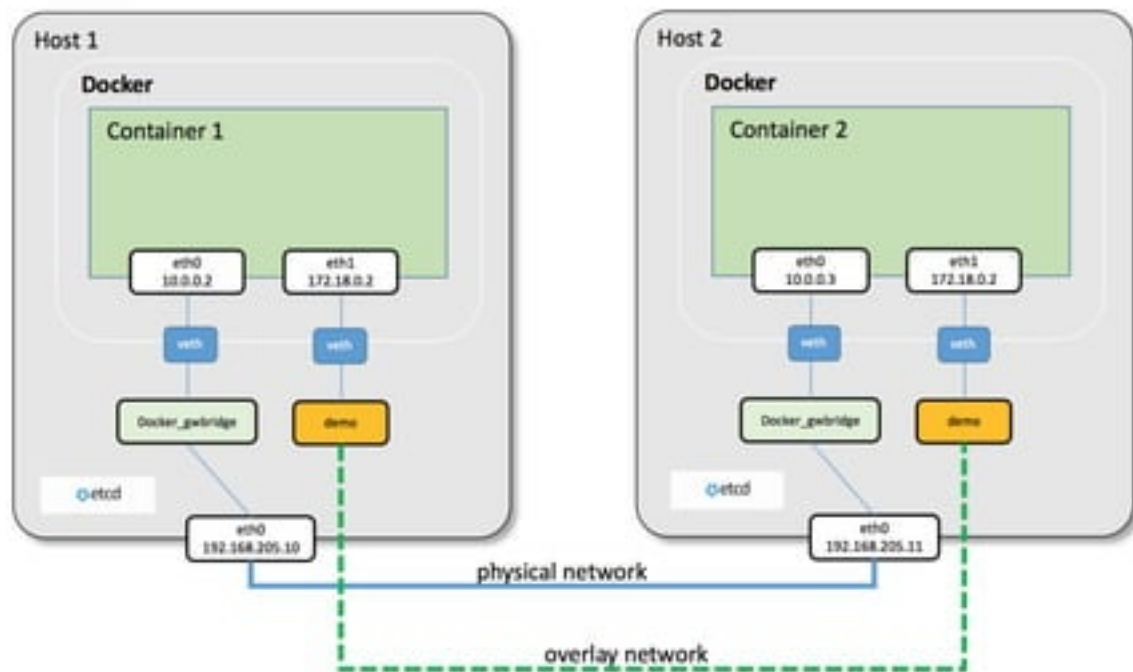
# Lab A: Overlay Multi-Host Networking with OVS

- <http://docker-k8s-lab.readthedocs.io/en/latest/docker/docker-ovs.html>



## Lab B: Multi-Host Overlay Networking with Etcd

- <http://docker-k8s-lab.readthedocs.io/en/latest/docker/docker-etcd.html>



## Docker 1.12 Networking Model Overview

Features	Engine 1.11(and prior)	Engine 1.12
Multi-host Networking & KV Store	External KV store	No External KV Store required (Swarm Mode)
MACVLAN	Experimental	Out-of Experimental
Secure Control Plane	Insecure (Plain-text)	Secure
Secure Data Plane	VXLAN was not encrypted by default(can be secured by --opt-secure)	Encrypted VXLAN traffic( makes use of swarm certificates and key exchange)
Load Balancing	Load Balancing was featured under 1.10, based on only DNS RR	Virtual IP Load-Balancing & DNS RR both supported(LB using IPVS)
Service Discovery	Available under 1.10 but based on external service discovery backend	Service discovery now integrated into Docker Engine, Virtual IP for VIP Load-Balancing support.
Swarm Mode	Not Available	Newly Introduced ( Optional Feature)
Routing Mesh	Not Available	Newly Introduced

# Contents

---

- The Metrix from Hell
- What is Container/Docker?
- Docker Basic Usage
- Docker Networking Deep Dive
- **Docker Compose**
- Docker Swarm
- Ecosystem and Standardization

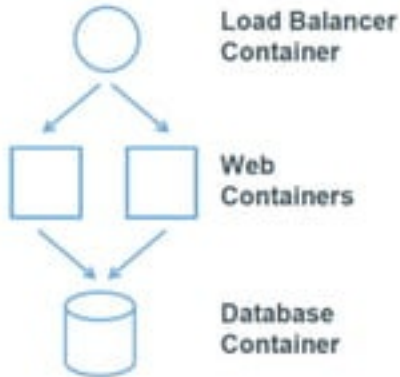


# Docker Compose Overview

- Compose is a tool for defining and running multi-container Docker applications. With Compose, you use a Compose file to configure your application's services. Then, using a single command, you create and start all the services from your configuration.
- Defined in yaml

```
web:  
  build: .  
  links:  
    - db  
  ports:  
    - "8000:8000"  
db:  
  image: postgres
```

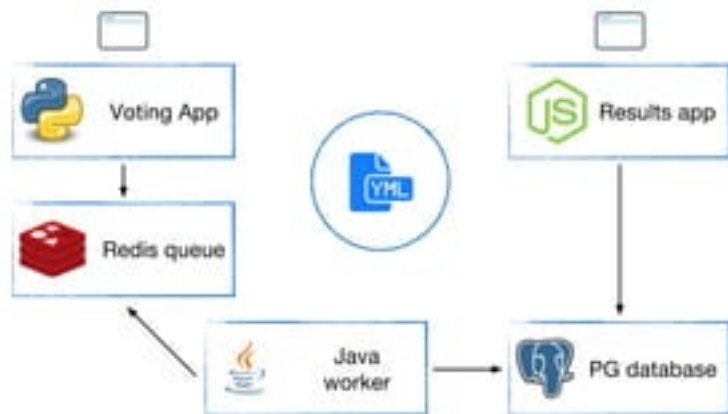
## Docker Compose



# Demo: Docker Compose

---

- Example Voting App
  - <https://github.com/DaoCloud/example-voting-app>
- Guide:
  - <http://docker-k8s-lab.readthedocs.io/en/latest/docker/docker-compose.html>





# Contents

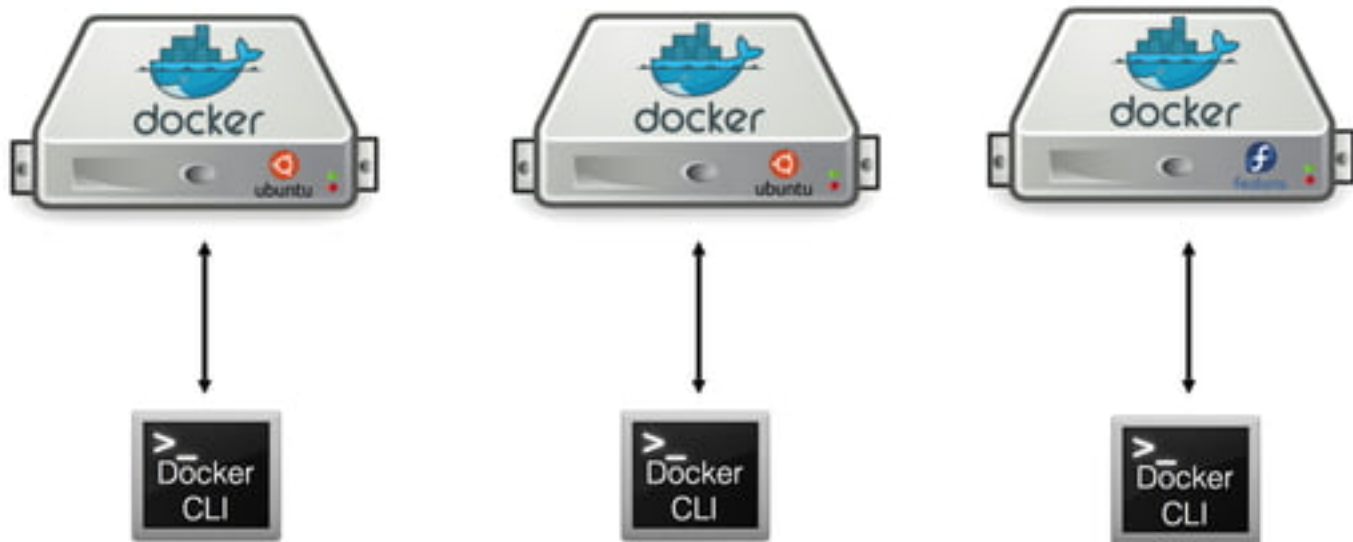
---

- The Metrix from Hell
- What is Container/Docker?
- Docker Basic Usage
- Docker Networking Deep Dive
- Docker Compose
- **Docker Swarm**
- Ecosystem and Standardization

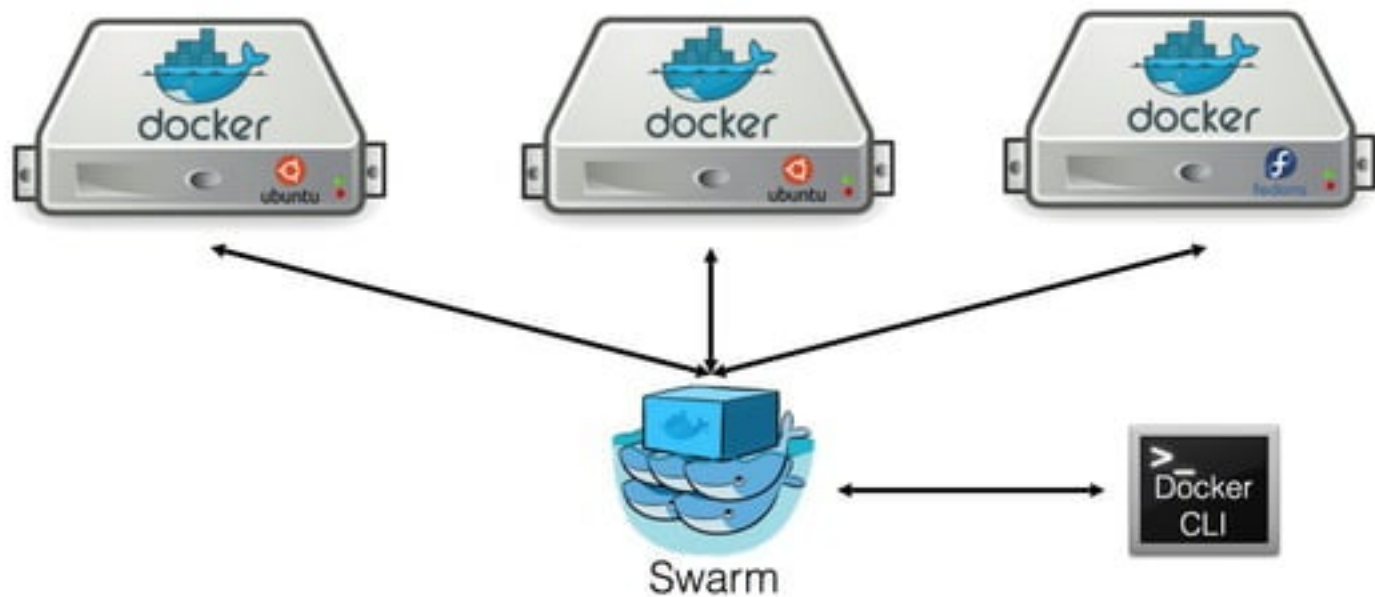


## Before Docker Swarm

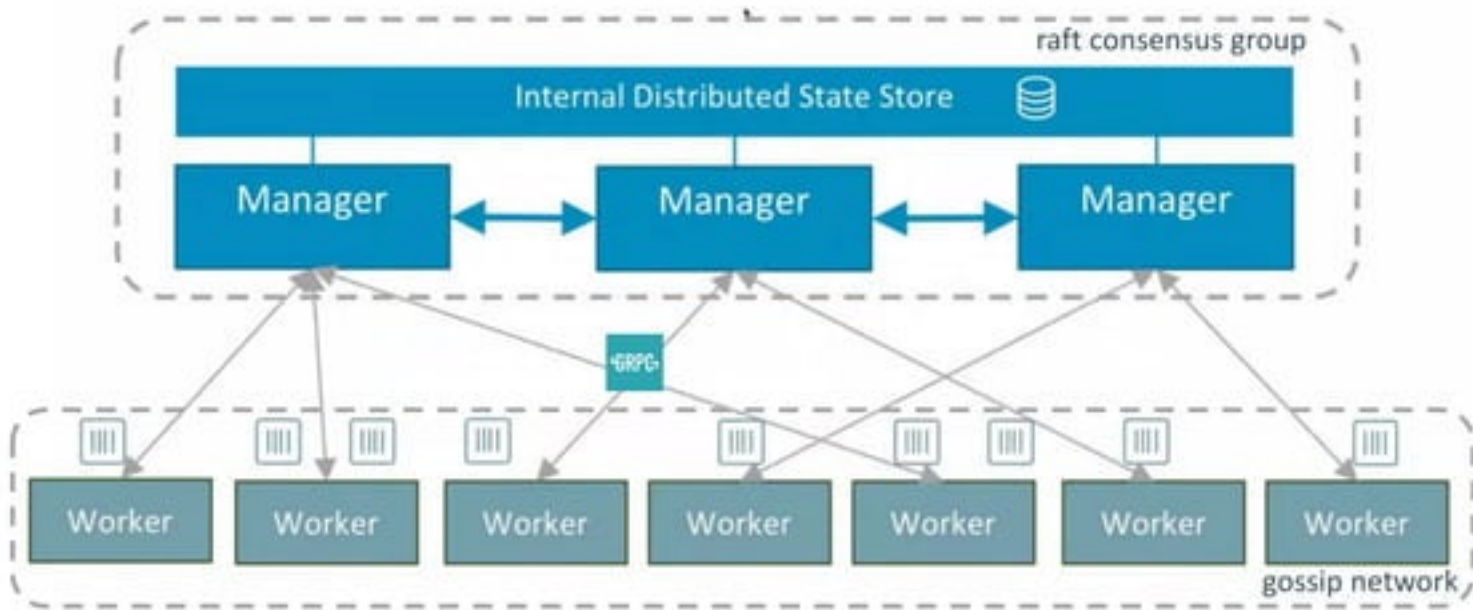
---



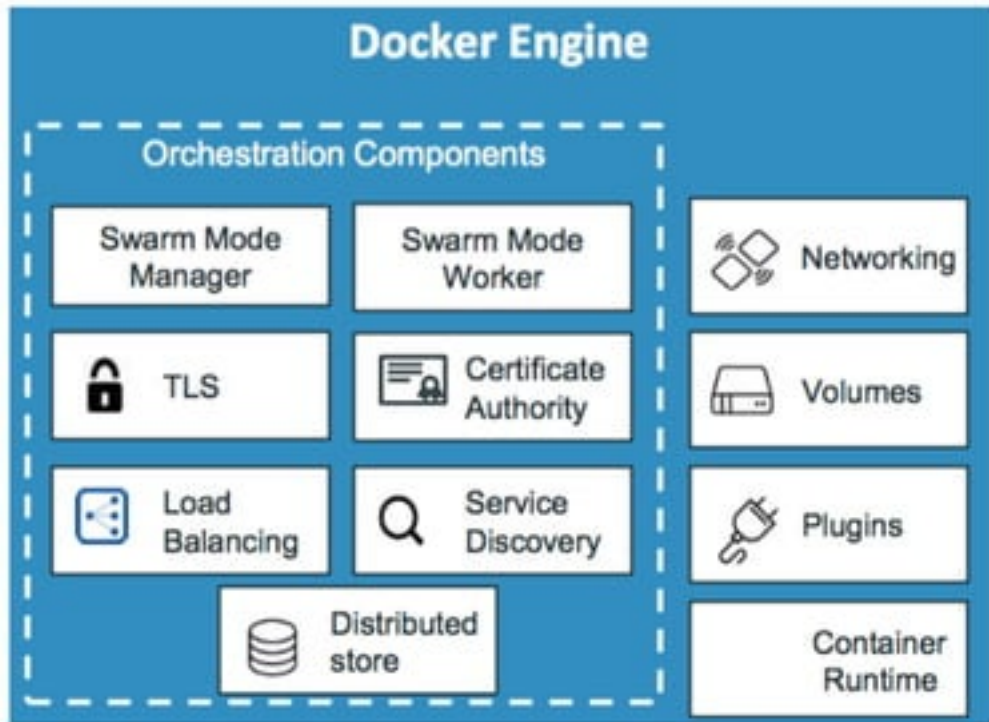
## With Docker Swarm



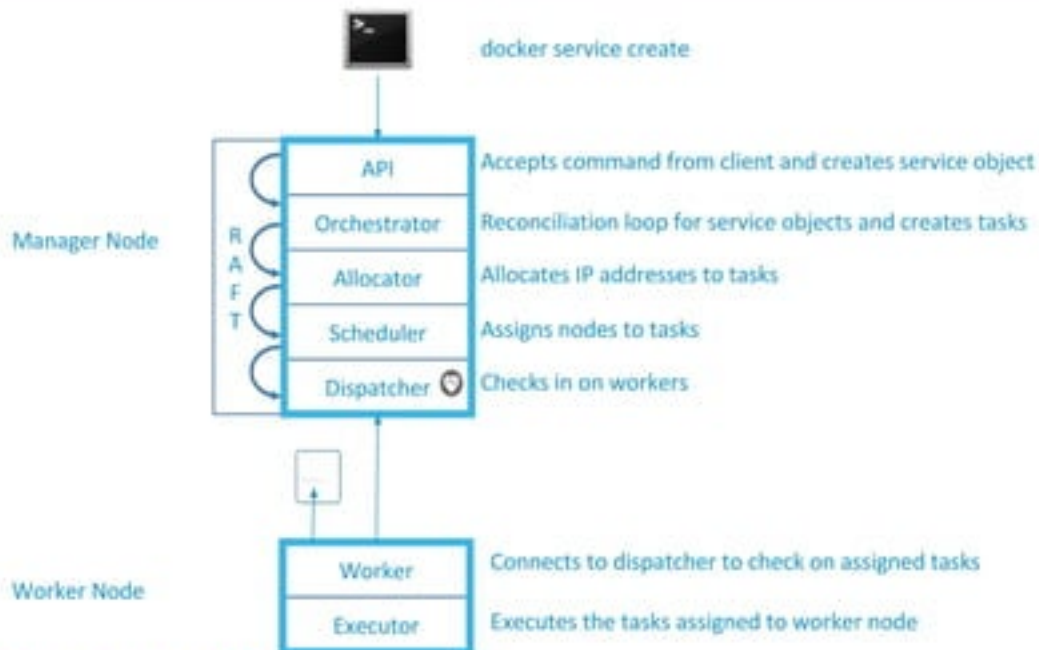
# Swarm mode cluster architecture



# Docker Orchestration Components



# Create Docker Service



<http://collabnix.com/archives/1445>

[https://www.youtube.com/watch?v=\\_F6PSP-qhdA](https://www.youtube.com/watch?v=_F6PSP-qhdA)

## Demo: Docker Swarm

---

- Docker Swarm with Load Balancing and Scaling
  - <http://docker-k8s-lab.readthedocs.io/en/latest/docker/docker-swarm-lb-scale.html>
- Multi-Services issue
  - <https://github.com/docker/compose/issues/3656>

# Contents

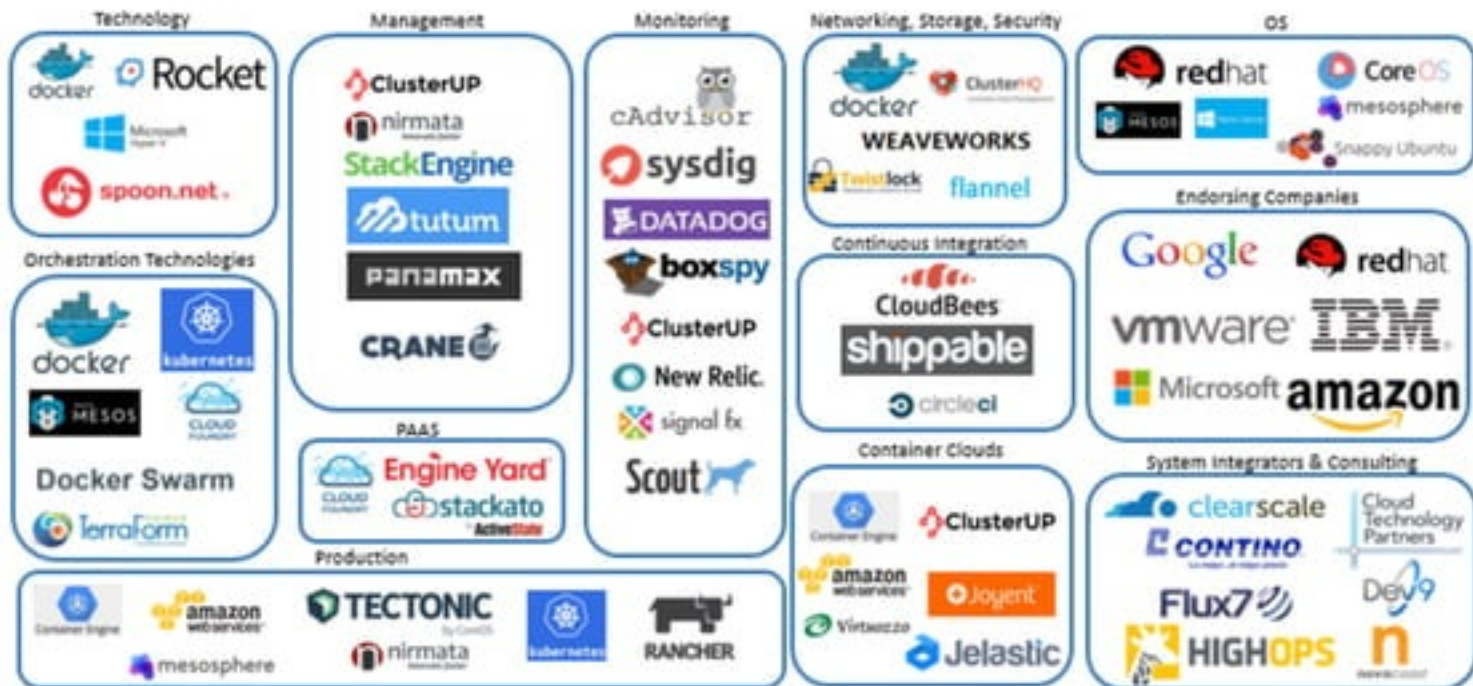
---

- The Metrix from Hell
- What is Container/Docker?
- Docker Basic Usage
- Docker Networking Deep Dive
- Docker Compose
- Docker Swarm
- **Ecosystem and Standardization**





# Container Ecosystem



# Open Container Initiative (OCI)

---

- A **Linux Foundation** Collaborative Project
- Free from control by any particular vendor's specific cloud stack or ecosystem
- Includes:
  - Container runtime specification -> runc
  - Image format specification



<http://www.slideshare.net/PhilEstes/runc-the-little-engine-that-could-run-docker-containers>  
<https://www.opencontainers.org/>  
<https://github.com/opencontainers>



# Container Networking

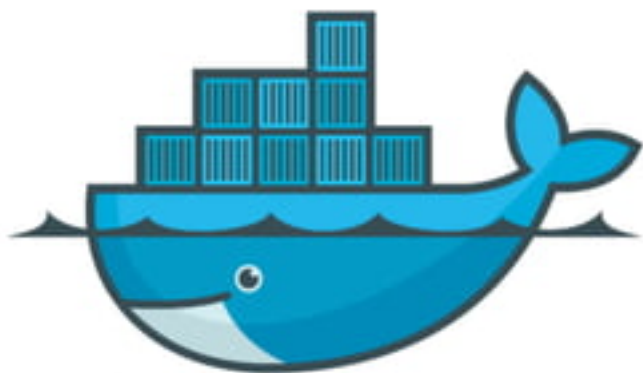
---

- Container Network Model (CNM) by Docker
  - Adopted by Docker libnetwork, Cisco Contiv, Kuryr, OVN, Project Calico, Vmware, Vwave
- Container Network Interface(CNI) by CoreOS
  - Adopted by kuberntes, Kurma, rkt, Apache Mesos, Cloud Foundry, Cisco Contiv, Project Calico and Weave.

<https://github.com/containernetworking>

<http://thenewstack.io/container-networking-landscape-cni-coreos-cnm-docker/>





docker  
?